

1. Fall 2026 entries:
 - a. 30610
 - b. I used COUNT(*) to count all the rows in the table WHERE the term column held “Fall 2026” exactly.
2. Percentage of international students:
 - a. 47.01%
 - b. I used COUNT(*) FILTER to isolate the us_or_international “International” count, and put in a NULLIF in the denominator COUNT to avoid division by zero if the database is empty. The decimal is then multiplied by 100 to make it a percentage, and ROUNDED to 2 decimal places.
3. Average G... metrics:
 - a. GPA: 3.78, GRE-Q: 165.4, GRE-V: 161.2, GRE-AW: 4.4
 - b. I used AVG to get the average of each G-metric, and ROUNDED GPA to 2 decimal places, and the rest to 1 decimal place. I also implemented a bit of additional data cleaning that I had not thought of originally. Each score has to fall BETWEEN its respective minimums AND maximums. These results are cast to a numeric type so that ROUND works.
4. Average American student GPA:
 - a. 3.79
 - b. I pulled the data FROM applicants WHERE the term was “Fall 2026” AND us_or_international = “American” and then just took the AVG of that data and ROUNDED to 2 decimal places, after converting to numeric data.
5. Fall 2026 acceptance rate:
 - a. 36.70%
 - b. This one was very similar in structure to question #2, except I used ILIKE ‘Accepted%’ to ignore the extra text inside the status indicator and make it case-insensitive.
6. Average GPA of Fall 2026 acceptances:
 - a. 3.76
 - b. I took data FROM applicants WHERE the term was “Fall 2026” AND their status included the word “Accepted” in the same way as problem #5. Then, I just took the AVG, and ROUNDED it.
7. JHU Comp. Sci. masters count:
 - a. 17
 - b. I COUNTed all applicants WHERE llm_generated_university = ‘Johns Hopkins University’ AND degree = ‘Masters’ AND used ILIKE again with % to catch anything inside llm_generated_program that contained “Computer Science”
8. 2026 acceptance entries for PhD students at specific universities:
 - a. 30
 - b. I COUNTed FROM the list of applicants WHERE the term contained “2026” AND status contained “Accepted” AND degree = “PhD”. I narrowed it further and avoided a

bunch of OR statements by adding AND llm_generated_university IN ((the list of all the schools)) AND checking that the llm_generated_program contained “computer science”

9. LLM fields vs. downloaded fields:

- a. 30 and 30
- b. This one checks the same data as the one above, except it uses the raw fields. I had to write a bunch of OR statements so that I could try to capture a few more of the of the possible names (like MIT vs. Massachusetts Institute of Technology)

10. Which 5 programs had the highest volume of applications?

- a. - Computer Science: 2470 applications - Economics: 2087 applications - Physics: 1968 applications - Philosophy: 1855 applications - English: 1557 applications
- b. To answer this question, I GROUPed BY llm_generated_program to get the total number of applications per program name. Then I ORDERed BY volume in DESCending order, with a LIMIT of 5.

11. Which months have the most results posted?

- a. - February: 17979 entries- March: 17325 entries - January: 5111 entries - April: 3225 entries - May: 249 entries - June: 30 entries
- b. To answer this, I had to make sure I stored date_added appropriately, as a SQL DATE object. I used PostgreSQL's TO_CHAR(date_added, 'Month') function to extract just the month string. Then I COUNTed them, and put them in DESCending order.